CSU33061 Artificial Intelligence I

Introduction

www.scss.tcd.ie/Tim.Fernando/AI

Key Phrases:

Can machines think?

- Turing test & ELIZA effect
- AI-complete (contra low hanging fruit)

Agent & environment

- Cognitive Revolution & Big Data

Levels of intelligence

Can machines think? (Turing 1950)

Turing test: can C tell A from B?

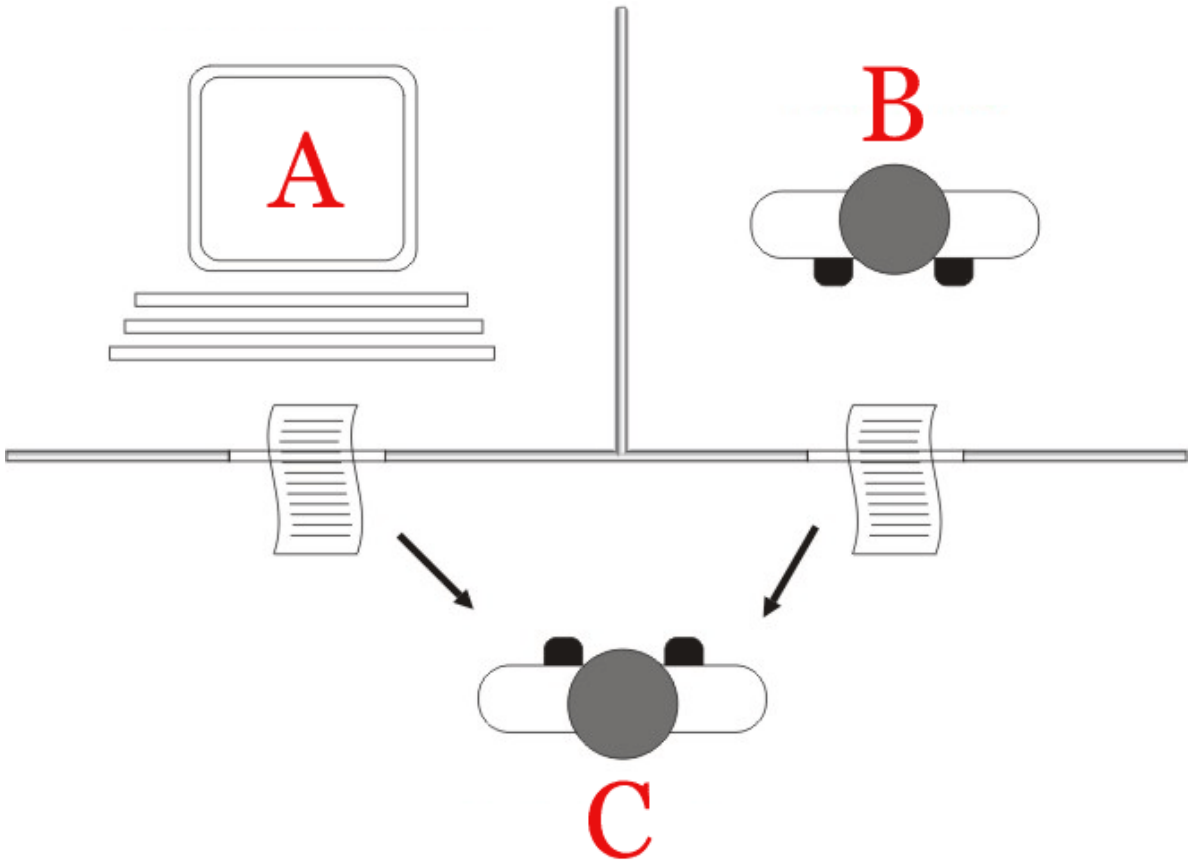

From Wikipedia, (Juan Alberto Sánchez Margallo)

Intelligence operationalized: subject to testing

Can machines think? (Turing 1950)

Turing test: can C tell A from B?

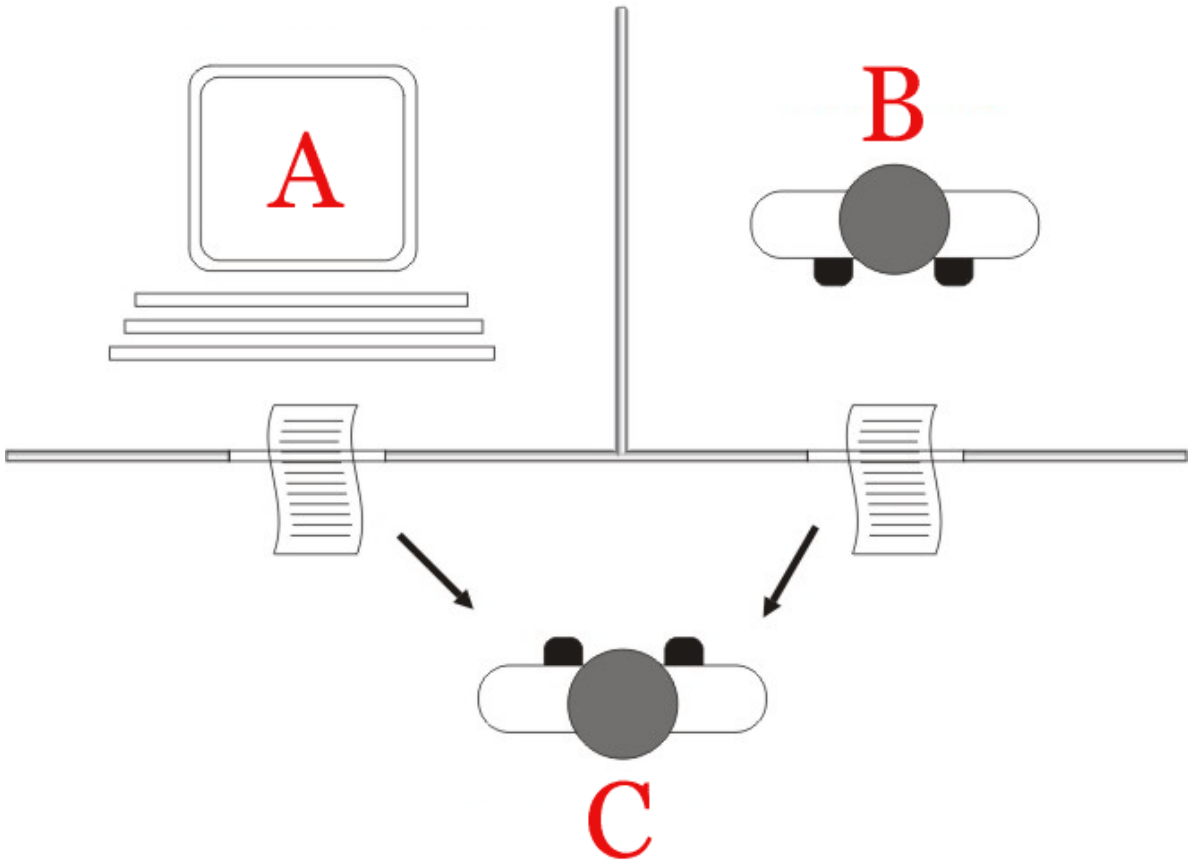

From Wikipedia, (Juan Alberto Sánchez Margallo)

Intelligence operationalized: subject to testing

. . . cheating?

ELIZA (Weizenbaum, 1964-66) & artful deception

- use pattern matching and substitution to fake understanding

ELIZA (Weizenbaum, 1964-66) & artful deception

- use pattern matching and substitution to fake understanding

ELIZA effect: humans are inclined to see computers as humans

e.g. when ATM says “thank you”

ELIZA (Weizenbaum, 1964-66) & artful deception

- use pattern matching and substitution to fake understanding

ELIZA effect: humans are inclined to see computers as humans

e.g. when ATM says "thank you"

An AI problem is **AI-complete** if any AI problem is mechanically reducible to it (i.e., it is at least as hard as any other).

E.g. Natural Language Understanding

The town councilors refused to give the demonstrators a permit because they feared violence.
Who feared violence?

T. Winograd

ELIZA (Weizenbaum, 1964-66) & artful deception

- use pattern matching and substitution to fake understanding

ELIZA effect: humans are inclined to see computers as humans

e.g. when ATM says "thank you"

An AI problem is **AI-complete** if any AI problem is mechanically reducible to it (i.e., it is at least as hard as any other).

E.g. Natural Language Understanding

The town councilors refused to give the demonstrators a permit because they advocated violence.
Who advocated violence?

T. Winograd

ELIZA (Weizenbaum, 1964-66) & artful deception

- use pattern matching and substitution to fake understanding

ELIZA effect: humans are inclined to see computers as humans

e.g. when ATM says "thank you"

An AI problem is **AI-complete** if any AI problem is mechanically reducible to it (i.e., it is at least as hard as any other).

E.g. Natural Language Understanding

The town councilors refused to give the demonstrators a permit because they advocated violence.
Who advocated violence?

T. Winograd

CAUTION: Programs may appear to work better than they do

Siri rage (Urban dictionary):

When you get enraged because Siri just doesn't get it.

Chinese room argument (Searle's thought experiment)

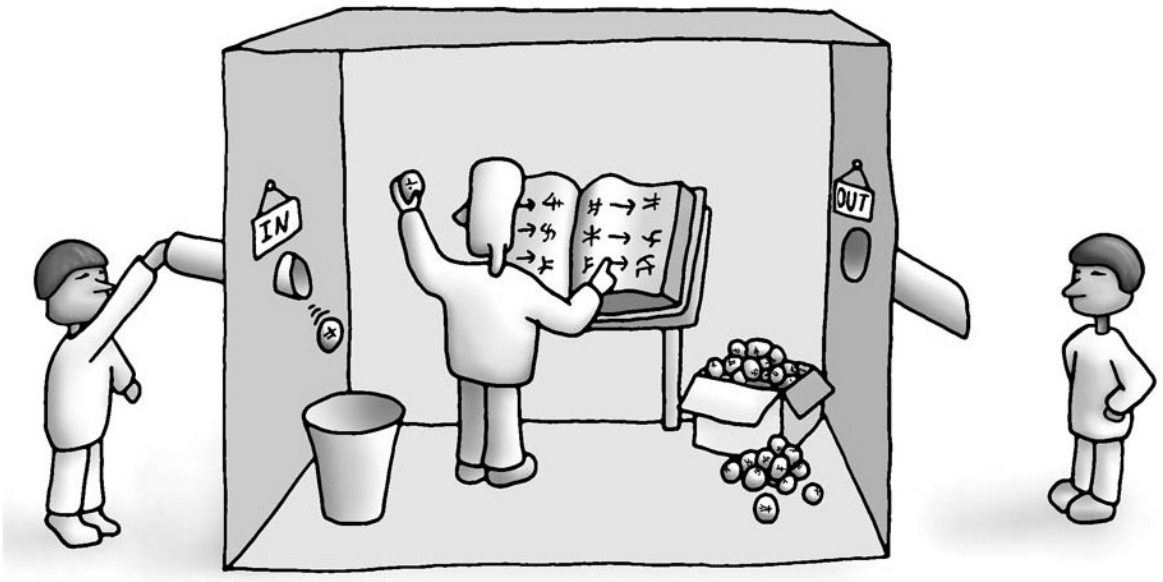

From http://america.pink/images/9/6/3/2/5/4/en/2-chinese-room.jpg

- a clerk can follow instructions for communicating in Chinese without understanding Chinese

Strong vs Weak AI

Chinese room argument (Searle's thought experiment)

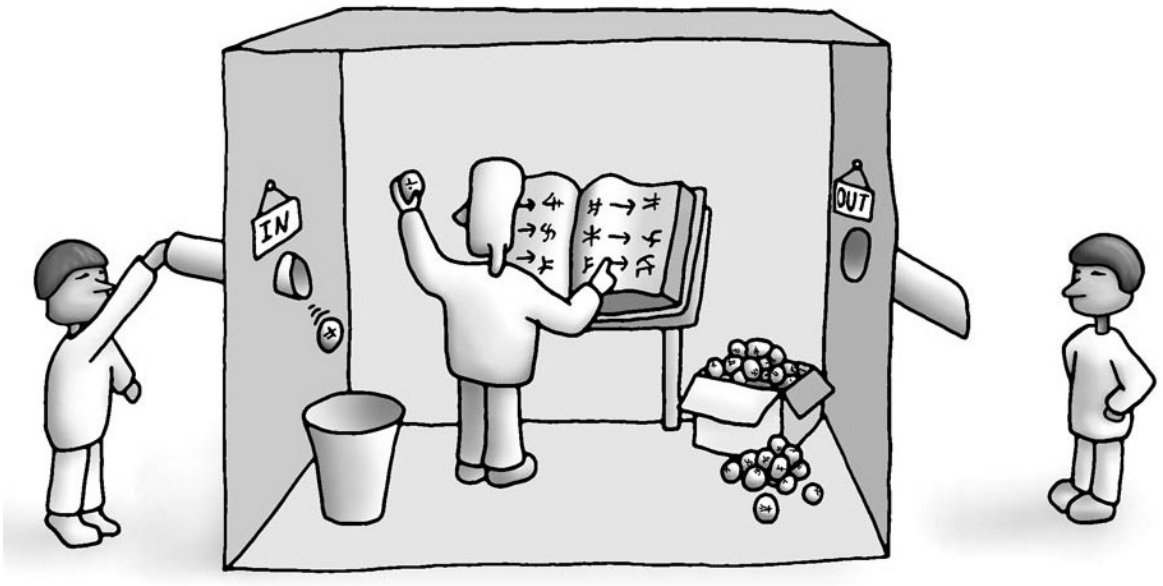

From http://america.pink/images/9/6/3/2/5/4/en/2-chinese-room.jpg

- a clerk can follow instructions for communicating in Chinese without understanding Chinese

mindless obedience $\neq$ mind from biological processes

Strong vs Weak AI

Chinese room argument (Searle's thought experiment)

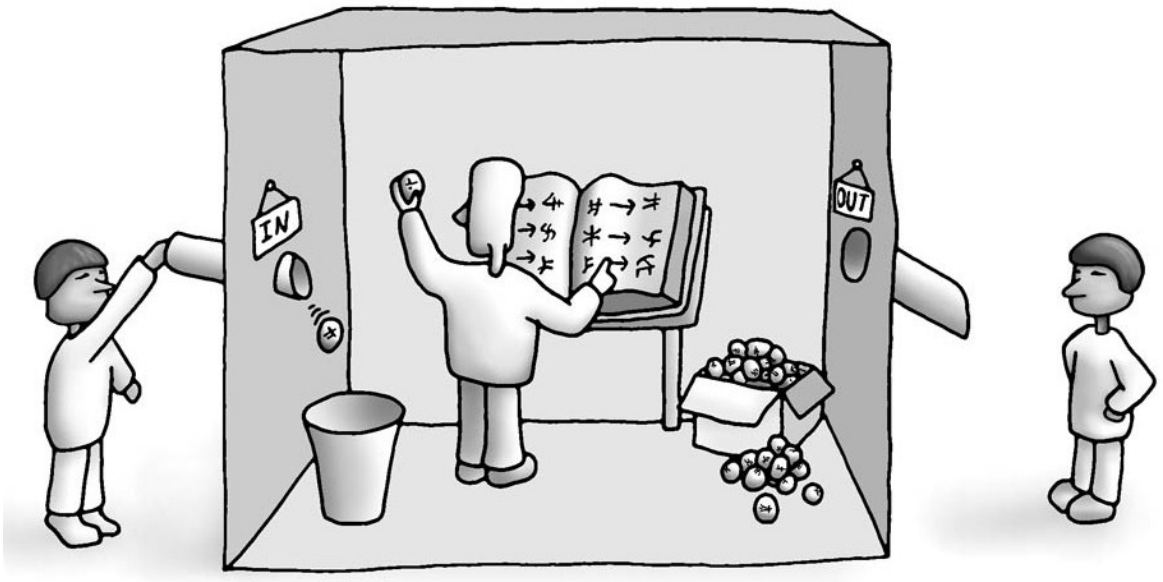

From http://america.pink/images/9/6/3/2/5/4/en/2-chinese-room.jpg

- a clerk can follow instructions for communicating in Chinese without understanding Chinese

mindless obedience $\neq$ mind from biological processes

Just do it $\rightsquigarrow$ Black Box, judged by its actions

Locating intelligence (black box)

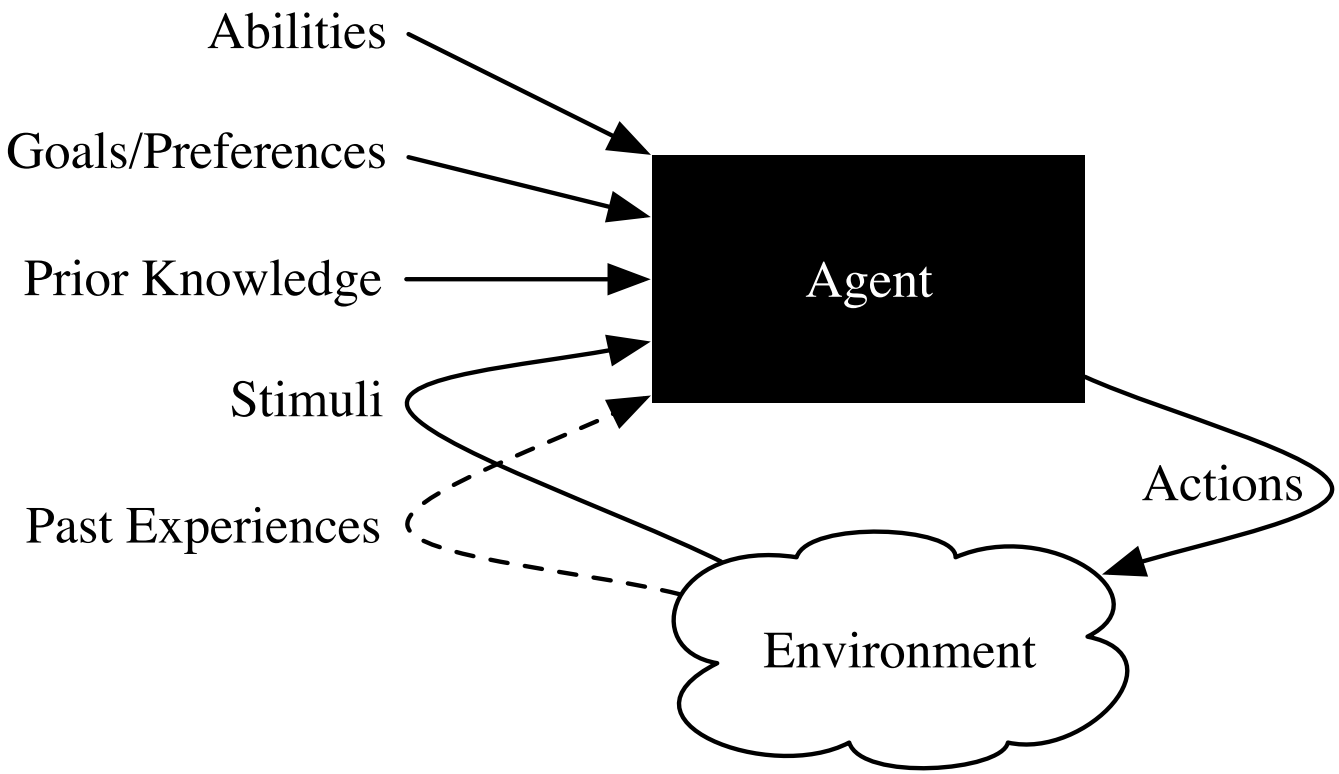

Poole & Mackworth

Intelligence: (abilities, goals, . . ., experience) $\mapsto$ action

Locating intelligence (black box)

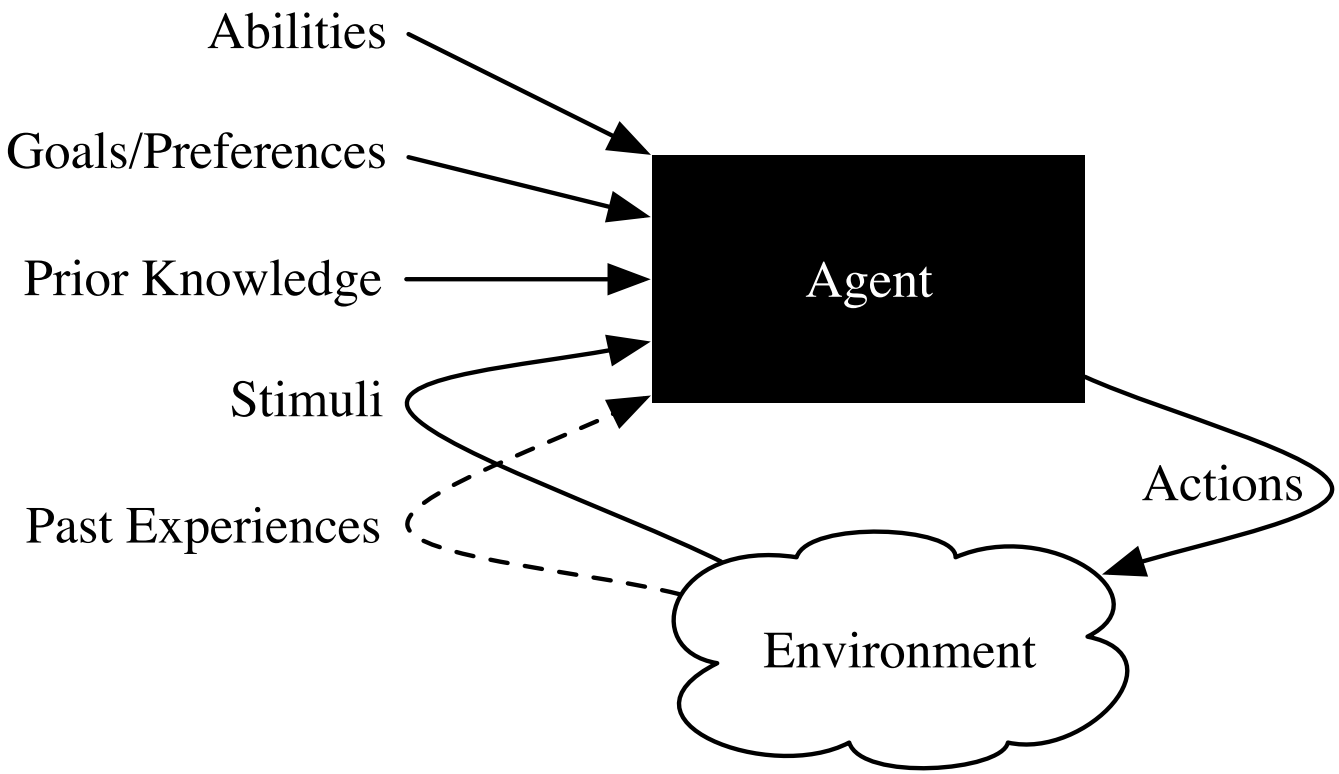

Poole & Mackworth

Intelligence: (abilities, goals, ..., experience) $\mapsto$ action

Turing test: what to say $\rightsquigarrow$ what to do

Between agent and environment

agent	environment
program	data
Cognitive Revolution	Big Data
hard-wired	experienced
rationalist	empiricist
nativist	behaviorist
innate	tabula rasa
nature	nurture

Turing machine &
specialized automaton

Between agent and environment

agent	environment
program	data
Cognitive Revolution	Big Data
hard-wired	experienced
rationalist	empiricist
nativist	behaviorist
innate	tabula rasa
nature	nurture

Turing machine &
specialized automaton

Learning (from environment)
trial & error: "data as oil"

Between agent and environment

agent	environment
program	data
Cognitive Revolution	Big Data
hard-wired	experienced
rationalist	empiricist
nativist	behaviorist
innate	tabula rasa
nature	nurture

Turing machine &
specialized automaton

Learning (from environment)
trial & error: "data as oil"

Moving target: changing agent & environment
e.g. change in state

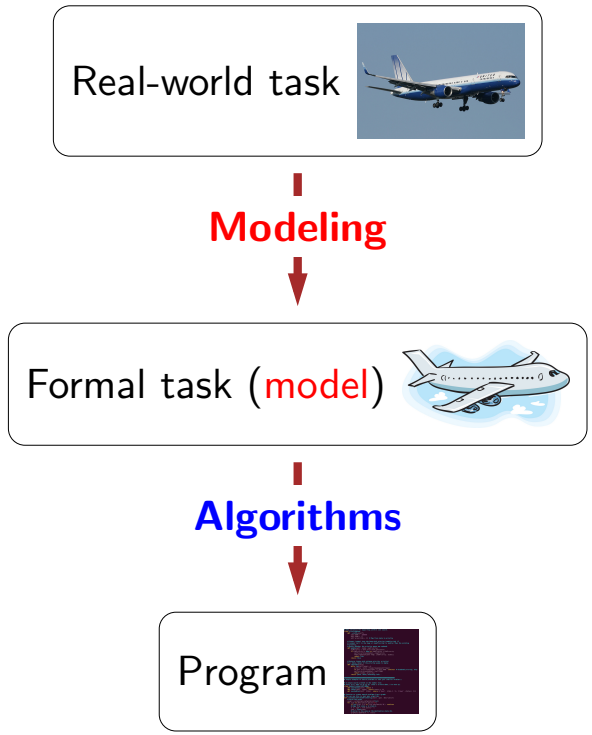

unstructured information $\rightsquigarrow$ actionable knowledge

Demis Hassabis

`www.theguardian.com/technology/2016/feb/16/demis-hassabis-artificial-intelligence-deepmind-alphago`

From Narrow to General AI

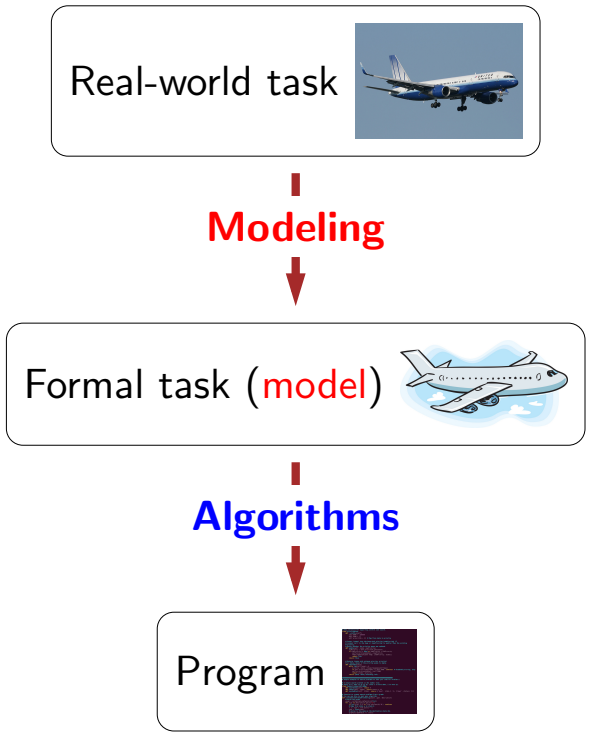

unstructured information $\rightsquigarrow$ actionable knowledge

Demis Hassabis

`www.theguardian.com/technology/2016/feb/16/demis-hassabis-artificial-intelligence-deepmind-alphago`

Autonomous = perform tasks in complex environments without constant user guidance

Adaptive = improve performance by learning from experience

From web.stanford.edu/class/cs221

Traditional approach

A spell checker:

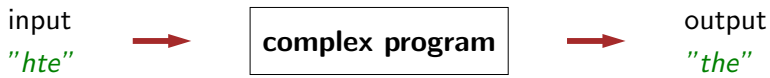

Problem: complexity becomes unwieldy

Machine learning approach

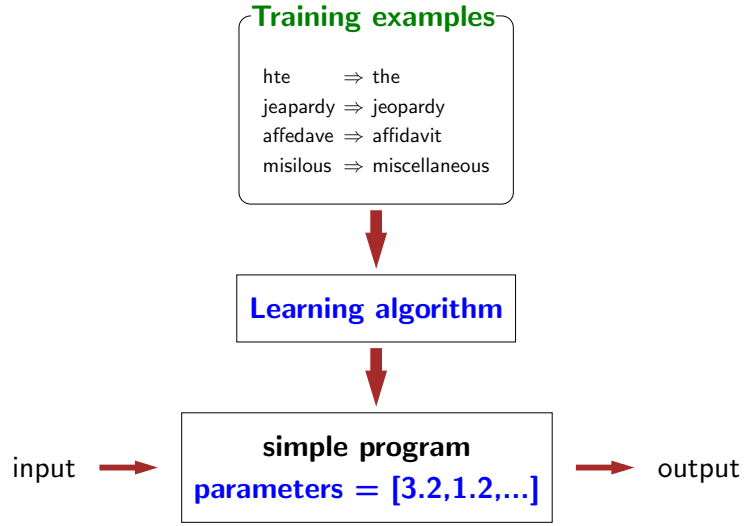

From web.stanford.edu/class/cs221

Traditional approach

A spell checker:

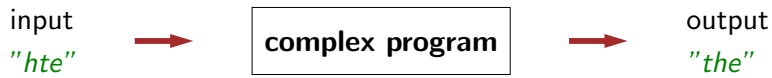

Problem: complexity becomes unwieldy

Machine learning approach

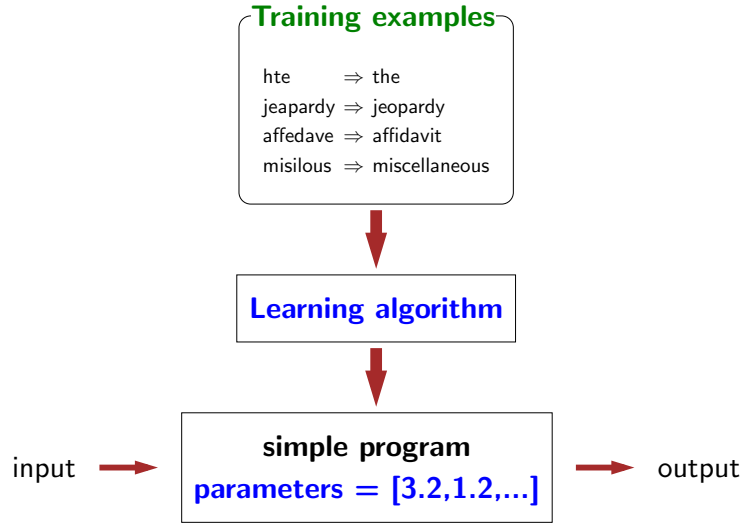

Search problems

Markov decision processes

Adversarial games

Constraint satisfaction problems

Bayesian networks

Reflex **States** **Variables** **Logic**

"Low-level intelligence" "High-level intelligence"

Machine learning

Back in Trinity

Undergraduate ML modules

- CSU44061 Machine Learning
  Semester 1 (5 ECTS)

- CSU44062 Advanced Computational Linguistics
  Semester 1 (5 ECTS)
  unsupervised ML for natural language processing

Back in Trinity

Undergraduate ML modules

- CSU44061 Machine Learning
  Semester 1 (5 ECTS)

- CSU44062 Advanced Computational Linguistics
  Semester 1 (5 ECTS)
  unsupervised ML for natural language processing

CSU33061: a taste building on CSU34011 (Prolog)

- logic & agents as Turing machines
- search
- Q-learning & Markov decision processes
- Constraint satisfaction
- Bayesian & Markov networks